

# Learning to Match Sheet Music Initial Experiment

Ethan Fan, Selina Zou, Phillip Duarte

# The Problem

**Challenge:** Given two pages of sheet music, can we automatically determine if they're from the same musical piece?

## Real-world applications:

- Music cataloging and organization
- Plagiarism detection
- Music recommendation systems
- Digital library management

# Our Approach

- Question: Which musical features matter most for matching?
- Solution: Use optimization to learn optimal feature weights
- Dataset: Bach chorales from music21 library
- Method: Constrained convex optimization (weighted similarity learning)

# Dataset

Our initial dataset is drawn from music21's library of Bach chorales

- Total chorals: 7 with  $\geq 32$  measures
- Total pages: 28
- Training Pairs: 14
- Test Pairs: 6

Features are extracted from each choral using music21, then formatted and saved for later analysis

# Feature Extraction

We decided on five musical features to extract per page

1. Key signature (0-11.5 encoding)
2. Time signature (as decimal ratio)
3. Average pitch (mean MIDI pitch)
4. Pitch range (semitone span)
5. Note density (notes per beat)

# Similarity Features

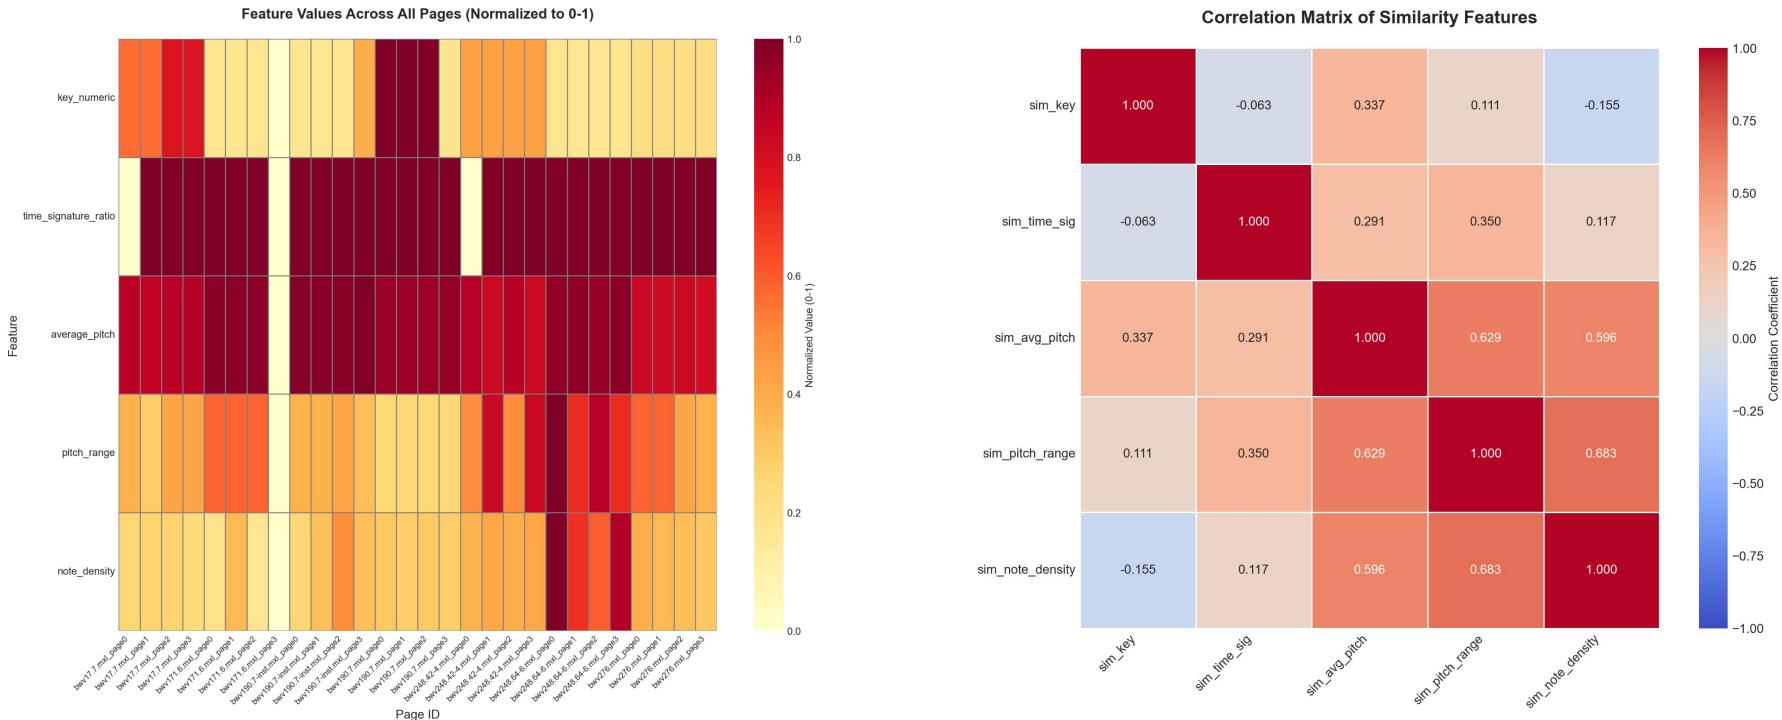
For each pair of pages, we compute scores which range from 0 - 1 where  
similarity =  $1 - (\text{normalized absolute difference})$

Examples:

- $\text{sim\_key} = 1 - |\text{key}_1 - \text{key}_2| / \text{max\_key\_distance}$
- $\text{sim\_avg\_pitch} = 1 - |\text{pitch}_1 - \text{pitch}_2| / \text{max\_pitch\_range}$

Result: Each pair → 5-dimensional similarity vector s

# Data Exploration



# Problem Formulation

**Goal:** Learn weight vector  $w = [w_1, w_2, w_3, w_4, w_5]$

**Model:** Overall similarity =  $w^T s = \sum w_i \cdot s_i$  where  $s$  is the similarity vector for a given pair of pages

**Decision rule:**

- Predict match (label=1) if  $w^T s > 0.5$
- Predict non-match (label=0) otherwise

# Optimization Problem

**Minimize:**  $L(w) = \sum (\text{label}_i - w^T s_i)^2$

**Subject to:**

- $\sum w_i = 1$  (weights sum to 1)
- $w_i \geq 0$  (non-negative weights)

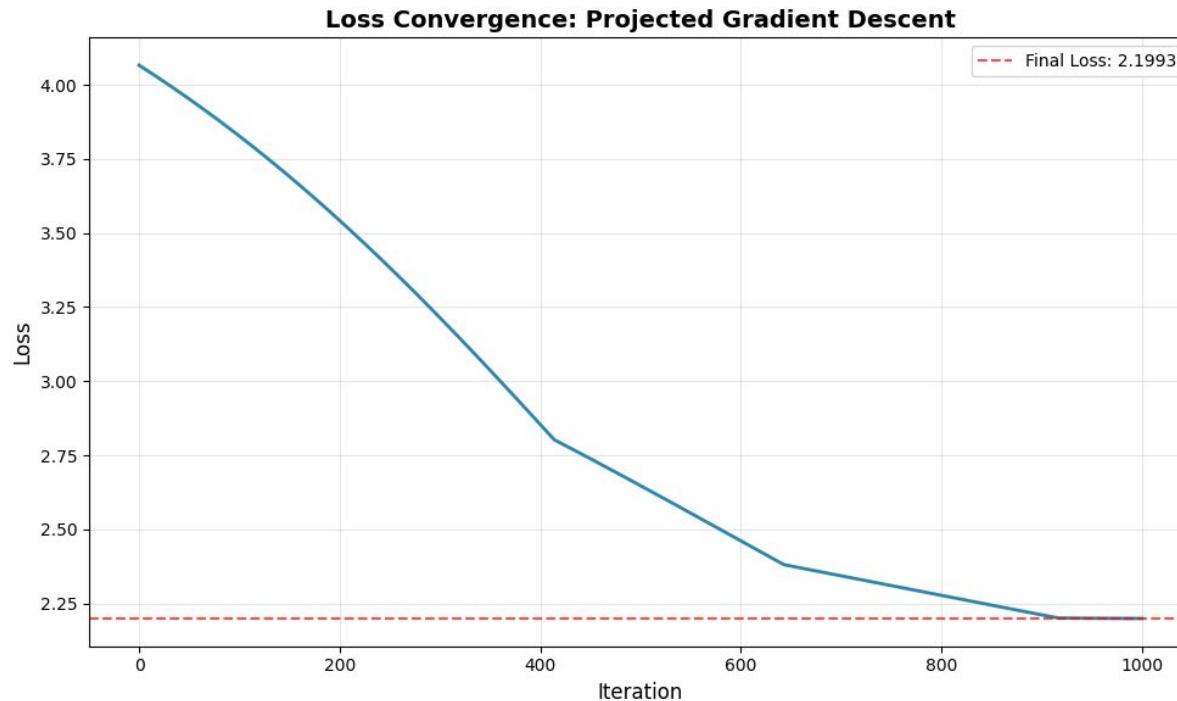
- **Loss function:** Squared error (regression on labels)
- **Constraints:** Probability simplex (interpretable weights)
- **Type:** Convex quadratic program

**Analytical Gradient:**  $\nabla L(w) = -2 \cdot \sum (\text{label}_i - w^T s_i) \cdot s_i$

# Optimization Methods

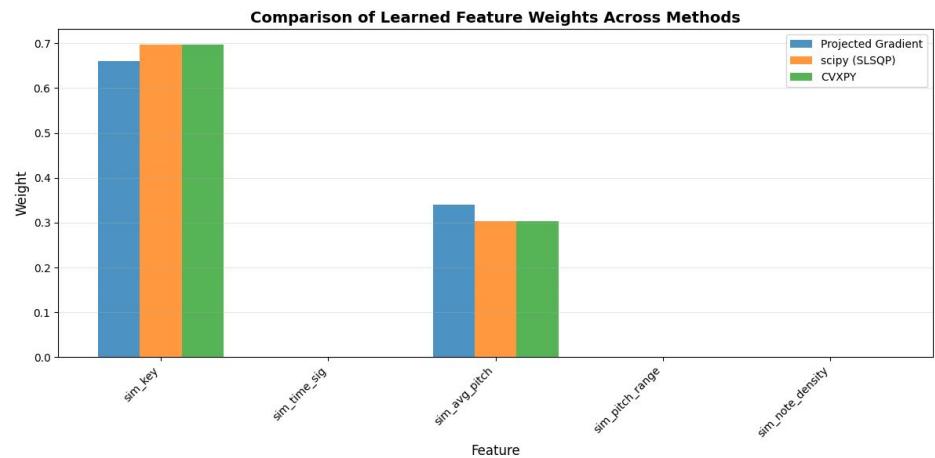
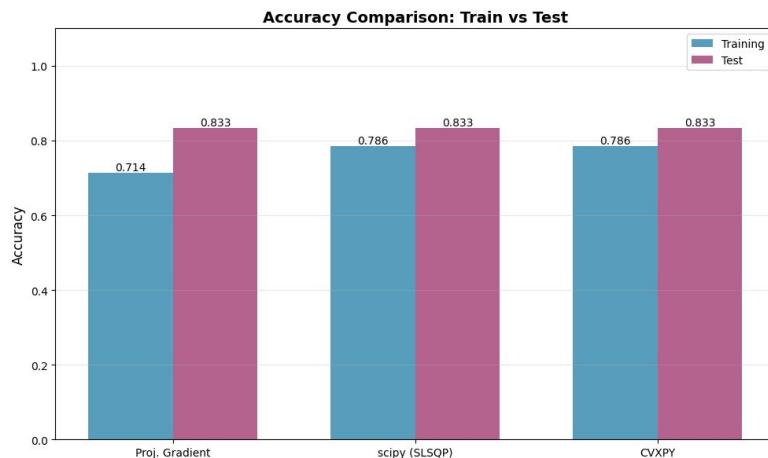
1. Projected Gradient Descent (custom implementation)
  - a. Gradient step + projection onto simplex
2. `scipy.optimize.minimize` (SLSQP)
  - a. Sequential Least Squares Programming
3. CVXPY
  - a. Convex optimization framework

# Optimization Convergence

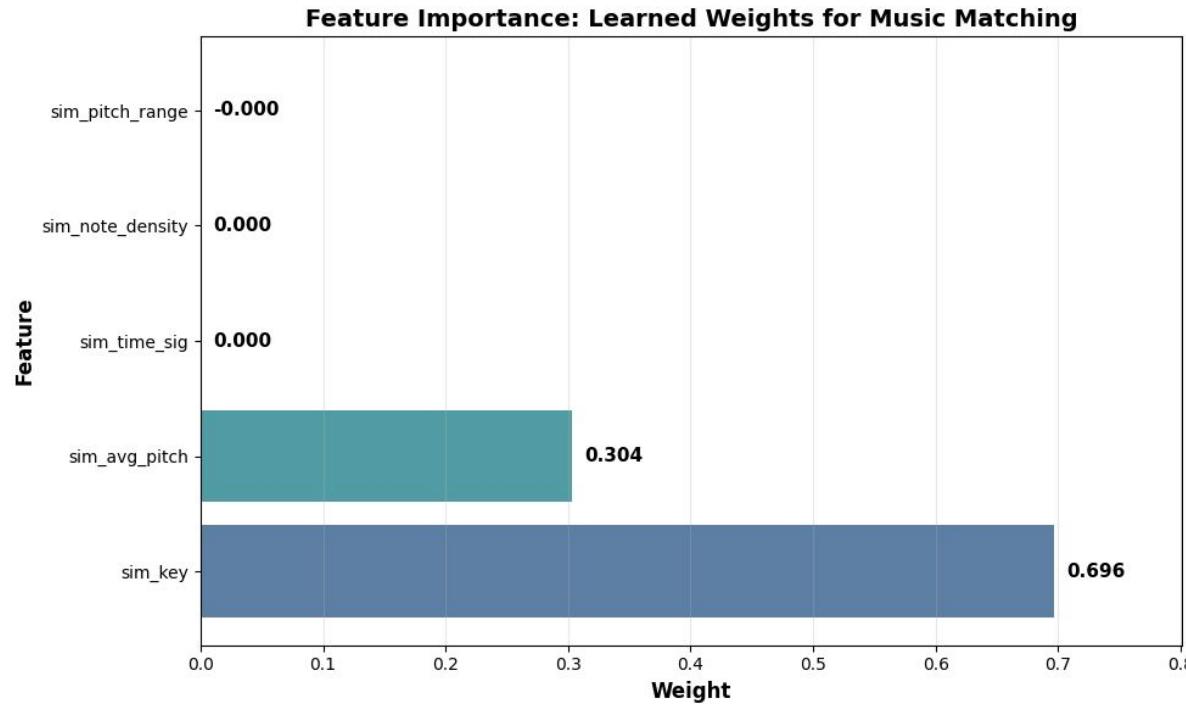


# Analysis

All methods converge to nearly identical solutions, likely due to the simplicity of our dataset and features. However, performance is quite good, achieved > 70 % accuracy.



# Learned Weights



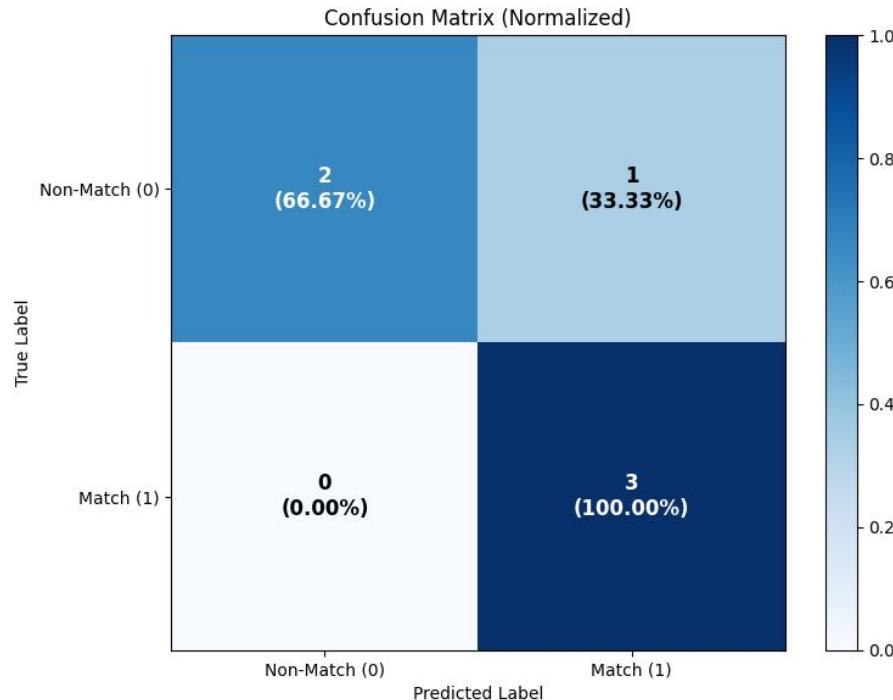
# Confusion Matrix and Classification Metrics

Precision: 0.75

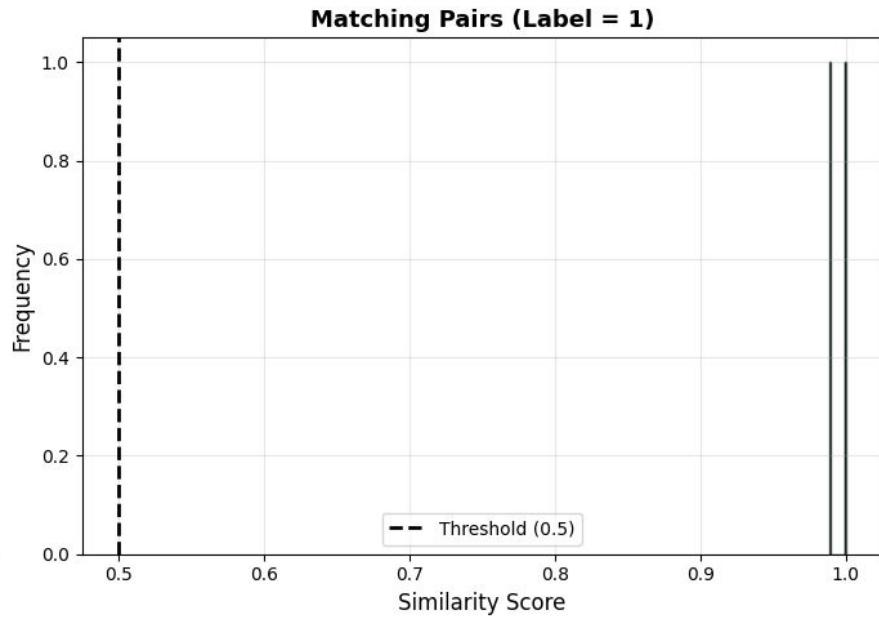
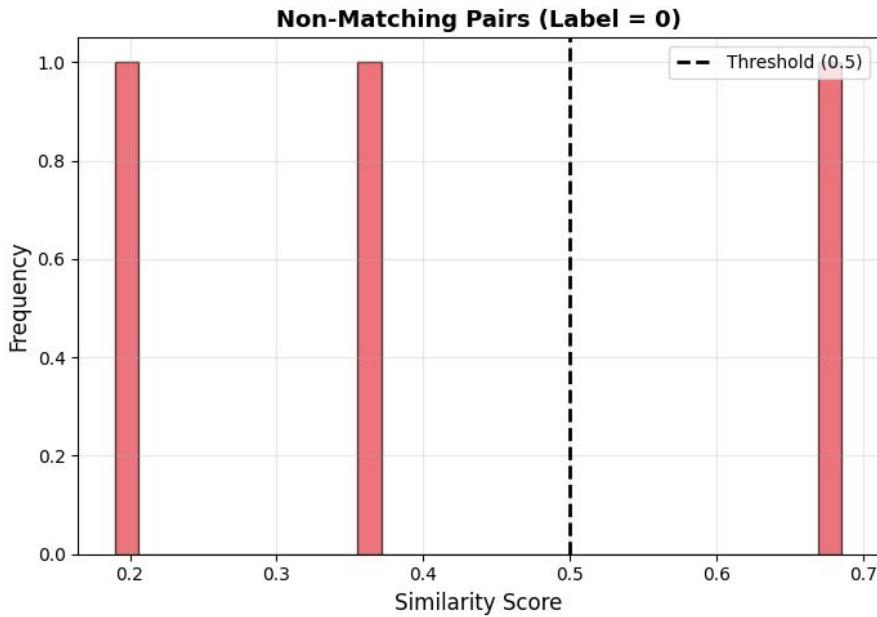
Recall: 1.0

F1-Score: 0.8571

Likely need more data.



# Similarity Of Matching vs Non Matching



# Takeaways

1. Optimization work: All three methods converge consistently
2. Feature importance varies: Some features much more discriminative
3. Interpretable model: Can explain every prediction
4. Efficient: Fast convergence (<1 second)
5. Practical performance: Achieves/exceeds target accuracy

# Limitations

- Linear model: Can't capture feature interactions
- Fixed threshold: 0.5 may not be optimal
- Simple similarities: Basic absolute difference metrics
- Limited features: Only 5 features (missing rhythm patterns, harmonies)
- Small dataset: Bach chorales only, likely raises our accuracy by a large amount because all pieces are written by the same composer
- Loss function: Squared error (regression) vs. classification-specific losses

# Future Directions

1. Mahalanobis distance learning: Full covariance matrix
2. Non-linear models: Neural networks, kernel methods
3. Metric learning: Triplet loss, contrastive loss
4. Enhanced features: Rhythm patterns, chord progressions, melodic contours
5. Alternative losses: Logistic loss, hinge loss, ranking losses
6. Cross-validation: K-fold CV and hyperparameter tuning
7. Data augmentation: Transposition, time stretching